

KITTI 3D Object Detection - Data Exploration and Visualization

This notebook provides comprehensive data exploration, visualization, and analysis for the KITTI 3D object detection dataset.

```
In [1]: import os
import cv2
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.decomposition import PCA
from sklearn.manifold import TSNE
from sklearn.cluster import KMeans, DBSCAN
from sklearn.preprocessing import StandardScaler
import torch
from config import cfg
from dataset import KITTIDataset, build_transforms

# Set style
sns.set_style('whitegrid')
plt.rcParams['figure.figsize'] = (12, 8)

print("Libraries imported successfully!")
```

Libraries imported successfully!

1. Dataset Overview

```
In [2]: # Load dataset
dataset = KITTIDataset(
    root=cfg.DATA_ROOT,
    split="training",
    transform=None
)

print(f"Total samples: {len(dataset)}")
print(f"Image directory: {dataset.image_dir}")
print(f"Label directory: {dataset.label_dir}")
```

Total samples: 7481
Image directory: C:\Users\94675\Desktop\group_project_pro\data\kitti\training\image_2
Label directory: C:\Users\94675\Desktop\group_project_pro\data\kitti\training\training\label_2

2. Data Distribution Analysis

```
In [3]: # Collect all 3D box parameters
all_boxes = []
for i in range(len(dataset)):
    _, target, _ = dataset[i]
```

```

    if target.sum() != 0: # Skip empty targets
        all_boxes.append(target.numpy())

all_boxes = np.array(all_boxes)
print(f"Valid samples with Car objects: {len(all_boxes)}")

# Create DataFrame
df = pd.DataFrame(all_boxes, columns=['x', 'y', 'z', 'length', 'width', 'height'])
df.head()

```

Valid samples with Car objects: 6684

Out[3]:

	x	y	z	length	width	height	rotation_y
0	-16.530001	2.39	58.490002	3.69	1.87	1.67	1.57
1	3.180000	2.27	34.380001	4.36	1.58	1.41	-1.58
2	1.000000	1.75	13.220000	4.15	1.73	1.57	1.62
3	-15.710000	2.16	38.259998	4.01	1.76	1.49	1.57
4	-2.720000	0.82	48.220001	3.62	1.56	1.48	-1.62

In [4]:

```

# Statistical summary
print("\n==== Statistical Summary ====")
print(df.describe())

```

==== Statistical Summary ====

	x	y	z	length	width	\
count	6684.000000	6684.000000	6684.000000	6684.000000	6684.000000	
mean	-2.248752	1.725507	24.496105	3.839042	1.652370	
std	6.829262	0.375542	17.427032	0.441941	0.099436	
min	-29.180000	-2.140000	-0.180000	2.330000	1.200000	
25%	-4.570000	1.600000	9.597500	3.560000	1.590000	
50%	-1.695000	1.720000	21.315000	3.830000	1.650000	
75%	1.150000	1.870000	35.830002	4.170000	1.720000	
max	39.860001	3.820000	81.660004	6.670000	2.040000	
	height	rotation_y				
count	6684.000000	6684.000000				
mean	1.530964	-0.111287				
std	0.138064	1.668206				
min	1.140000	-3.140000				
25%	1.440000	-1.580000				
50%	1.510000	-1.290000				
75%	1.610000	1.570000				
max	2.250000	3.140000				

In [5]:

```

# Distribution plots
fig, axes = plt.subplots(2, 4, figsize=(20, 10))
fig.suptitle('Distribution of 3D Box Parameters', fontsize=16, fontweight='bold')


columns = ['x', 'y', 'z', 'length', 'width', 'height', 'rotation_y']
for idx, col in enumerate(columns):
    row = idx // 4
    col_idx = idx % 4
    axes[row, col_idx].hist(df[col], bins=50, edgecolor='black', alpha=0.7)
    axes[row, col_idx].set_title(f'{col.capitalize()} Distribution')
    axes[row, col_idx].set_xlabel(col)
    axes[row, col_idx].set_ylabel('Frequency')
    axes[row, col_idx].grid(True, alpha=0.3)

```

```

# Remove empty subplot
fig.delaxes(axes[1, 3])
plt.tight_layout()
plt.savefig(os.path.join(cfg.VIS_DIR, 'parameter_distributions.png'), dpi=300, bbox_inches='tight')
plt.show()

```




3. Correlation Analysis

```

In [6]: # Correlation matrix
plt.figure(figsize=(10, 8))
correlation_matrix = df.corr()
sns.heatmap(correlation_matrix, annot=True, fmt='.2f', cmap='coolwarm',
            square=True, linewidths=1, cbar_kws={"shrink": 0.8})
plt.title('Correlation Matrix of 3D Box Parameters', fontsize=14, fontweight='bold')
plt.tight_layout()
plt.savefig(os.path.join(cfg.VIS_DIR, 'correlation_matrix.png'), dpi=300, bbox_inches='tight')
plt.show()

```



4. Dimensionality Reduction - PCA

```
In [7]: # Standardize data
scaler = StandardScaler()
data_scaled = scaler.fit_transform(df)

# Apply PCA
pca = PCA()
pca_result = pca.fit_transform(data_scaled)

# Explained variance
explained_variance = pca.explained_variance_ratio_
cumulative_variance = np.cumsum(explained_variance)

# Plot explained variance
fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(16, 6))

ax1.bar(range(1, len(explained_variance)+1), explained_variance, alpha=0.7, edgecolor='black')
ax1.set_xlabel('Principal Component')
ax1.set_ylabel('Explained Variance Ratio')
ax1.set_title('PCA - Explained Variance by Component')
ax1.grid(True, alpha=0.3)

ax2.plot(range(1, len(cumulative_variance)+1), cumulative_variance, 'bo-', linewidth=2)
ax2.axhline(y=0.95, color='r', linestyle='--', label='95% Variance')
ax2.set_xlabel('Number of Components')
```


```

ax2.set_ylabel('Cumulative Explained Variance')
ax2.set_title('PCA - Cumulative Explained Variance')
ax2.legend()
ax2.grid(True, alpha=0.3)

plt.tight_layout()
plt.savefig(os.path.join(cfg.VIS_DIR, 'pca_variance.png'), dpi=300, bbox_inches='tight')
plt.show()

print(f"\nVariance explained by first 3 components: {cumulative_variance[2]:.2%}")

```




Variance explained by first 3 components: 63.21%

```

In [8]: # 2D PCA visualization
plt.figure(figsize=(10, 8))
scatter = plt.scatter(pca_result[:, 0], pca_result[:, 1],
                      c=df['z'], cmap='viridis', alpha=0.6, edgecolors='black', linewidth=1)
plt.colorbar(scatter, label='Depth (z)')
plt.xlabel(f'PC1 ({explained_variance[0]:.1%} variance)')
plt.ylabel(f'PC2 ({explained_variance[1]:.1%} variance)')
plt.title('PCA - First Two Principal Components (colored by depth)', fontweight='bold')
plt.grid(True, alpha=0.3)
plt.tight_layout()
plt.savefig(os.path.join(cfg.VIS_DIR, 'pca_2d.png'), dpi=300, bbox_inches='tight')
plt.show()

```




5. Dimensionality Reduction - t-SNE

```
In [10]: # Apply t-SNE (on subset for speed)
import warnings
warnings.filterwarnings('ignore', category=FutureWarning)

subset_size = min(1000, len(data_scaled))
indices = np.random.choice(len(data_scaled), subset_size, replace=False)
data_subset = data_scaled[indices]

tsne = TSNE(n_components=2, random_state=42, perplexity=30, max_iter=1000) # 改
tsne_result = tsne.fit_transform(data_subset)

# Visualize t-SNE
plt.figure(figsize=(10, 8))
scatter = plt.scatter(tsne_result[:, 0], tsne_result[:, 1],
                      c=df.iloc[indices]['z'], cmap='plasma', alpha=0.6,
                      edgecolors='black', linewidth=0.5)
plt.colorbar(scatter, label='Depth (z)')
plt.xlabel('t-SNE Component 1')
plt.ylabel('t-SNE Component 2')
plt.title('t-SNE Visualization (colored by depth)', fontweight='bold')
plt.grid(True, alpha=0.3)
plt.tight_layout()
plt.savefig(os.path.join(cfg.VIS_DIR, 'tsne_2d.png'), dpi=300, bbox_inches='tight')
plt.show()
```




6. Clustering Analysis - K-Means

```
In [11]: # Elbow method for optimal K
inertias = []
K_range = range(2, 11)

for k in K_range:
    kmeans = KMeans(n_clusters=k, random_state=42, n_init=10)
    kmeans.fit(data_scaled)
    inertias.append(kmeans.inertia_)


# Plot elbow curve
plt.figure(figsize=(10, 6))
plt.plot(K_range, inertias, 'bo-', linewidth=2, markersize=8)
plt.xlabel('Number of Clusters (K)')
plt.ylabel('Inertia (Within-cluster sum of squares)')
plt.title('K-Means Elbow Method', fontweight='bold')
plt.grid(True, alpha=0.3)
plt.tight_layout()
plt.savefig(os.path.join(cfg.VIS_DIR, 'kmeans_elbow.png'), dpi=300, bbox_inches=
plt.show()
```



```
In [12]: # Apply K-Means with optimal K
optimal_k = 4
kmeans = KMeans(n_clusters=optimal_k, random_state=42, n_init=10)
clusters = kmeans.fit_predict(data_scaled)

# Visualize clusters in PCA space
plt.figure(figsize=(10, 8))
scatter = plt.scatter(pca_result[:, 0], pca_result[:, 1],
                      c=clusters, cmap='tab10', alpha=0.6,
                      edgecolors='black', linewidth=0.5)
plt.colorbar(scatter, label='Cluster')
plt.xlabel(f'PC1 ({explained_variance[0]:.1%} variance)')
plt.ylabel(f'PC2 ({explained_variance[1]:.1%} variance)')
plt.title(f'K-Means Clustering (K={optimal_k}) in PCA Space', fontweight='bold')
plt.grid(True, alpha=0.3)
plt.tight_layout()
plt.savefig(os.path.join(cfg.VIS_DIR, 'kmeans_clusters.png'), dpi=300, bbox_inches='tight')
plt.show()

# Cluster statistics
df['cluster'] = clusters
print("\n== Cluster Statistics ==")
print(df.groupby('cluster').mean())
```



==== Cluster Statistics ====

cluster	x	y	z	length	width	height	\
0	-12.222016	2.192860	39.915867	4.039538	1.728502	1.528716	
1	-4.235364	1.764063	15.095884	3.935024	1.615433	1.491928	
2	1.629682	1.345551	23.114407	4.214312	1.745169	1.690996	
3	0.760950	1.682082	26.285326	3.598885	1.623520	1.505510	

cluster	rotation_y
0	1.180338
1	1.645311
2	-0.314619
3	-1.547019

7. Clustering Analysis - DBSCAN

```
In [13]: # Apply DBSCAN
dbscan = DBSCAN(eps=0.5, min_samples=5)
dbscan_clusters = dbscan.fit_predict(data_scaled)

n_clusters = len(set(dbscan_clusters)) - (1 if -1 in dbscan_clusters else 0)
n_noise = list(dbscan_clusters).count(-1)

print(f"Number of clusters: {n_clusters}")
print(f"Number of noise points: {n_noise}")

# Visualize DBSCAN clusters
plt.figure(figsize=(10, 8))
```

```

scatter = plt.scatter(pca_result[:, 0], pca_result[:, 1],
                      c=dbSCAN_clusters, cmap='tab10', alpha=0.6,
                      edgecolors='black', linewidth=0.5)
plt.colorbar(scatter, label='Cluster (-1 = noise)')
plt.xlabel(f'PC1 ({explained_variance[0]:.1%} variance)')
plt.ylabel(f'PC2 ({explained_variance[1]:.1%} variance)')
plt.title(f'DBSCAN Clustering in PCA Space', fontweight='bold')
plt.grid(True, alpha=0.3)
plt.tight_layout()
plt.savefig(os.path.join(cfg.VIS_DIR, 'dbSCAN_clusters.png'), dpi=300, bbox_inches='tight')
plt.show()

```

Number of clusters: 263

Number of noise points: 835



8. Feature Engineering Insights

```

In [15]: # Create additional features
df['volume'] = df['length'] * df['width'] * df['height']
df['aspect_ratio'] = df['length'] / df['width']
df['distance_3d'] = np.sqrt(df['x']**2 + df['y']**2 + df['z']**2)

# Visualize new features
fig, axes = plt.subplots(1, 3, figsize=(18, 5))
fig.suptitle('Engineered Features Distribution', fontsize=14, fontweight='bold')

axes[0].hist(df['volume'], bins=50, edgecolor='black', alpha=0.7, color='skyblue'
axes[0].set_title('Volume Distribution')
axes[0].set_xlabel('Volume (m³)')

```

```


axes[0].set_ylabel('Frequency')
axes[0].grid(True, alpha=0.3)

axes[1].hist(df['aspect_ratio'], bins=50, edgecolor='black', alpha=0.7, color='lightcoral')
axes[1].set_title('Aspect Ratio Distribution')
axes[1].set_xlabel('Length / Width')
axes[1].set_ylabel('Frequency')
axes[1].grid(True, alpha=0.3)

axes[2].hist(df['distance_3d'], bins=50, edgecolor='black', alpha=0.7, color='lightgreen')
axes[2].set_title('3D Distance Distribution')
axes[2].set_xlabel('Distance (m)')
axes[2].set_ylabel('Frequency')
axes[2].grid(True, alpha=0.3)

plt.tight_layout()
plt.savefig(os.path.join(cfg.VIS_DIR, 'engineered_features.png'), dpi=300, bbox_inches='tight')
plt.show()

```



10. Summary Statistics

```

In [16]: # Generate comprehensive summary
summary = {
    'Total Samples': len(dataset),
    'Valid Car Objects': len(all_boxes),
    'Average Depth (z)': df['z'].mean(),
    'Average Length': df['length'].mean(),
    'Average Width': df['width'].mean(),
    'Average Height': df['height'].mean(),
    'PCA Components for 95% Variance': np.argmax(cumulative_variance >= 0.95) + 1,
    'Optimal K-Means Clusters': optimal_k,
    'DBSCAN Clusters': n_clusters,
    'DBSCAN Noise Points': n_noise
}

print("\n== Dataset Summary ==")
for key, value in summary.items():
    if isinstance(value, float):
        print(f"{key}: {value:.2f}")
    else:
        print(f"{key}: {value}")

# Save summary to file
with open(os.path.join(cfg.LOG_DIR, 'data_exploration_summary.txt'), 'w') as f:
    f.write("== Dataset Exploration Summary ==\n\n")
    for key, value in summary.items():
        if isinstance(value, float):
            f.write(f"{key}: {value:.2f}\n")

```

```
        else:
            f.write(f"\n{key}: {value}\n")

print("\nSummary saved to:", os.path.join(cfg.LOG_DIR, 'data_exploration_summary')

==== Dataset Summary ====
Total Samples: 7481
Valid Car Objects: 6684
Average Depth (z): 24.496105194091797
Average Length: 3.8390424251556396
Average Width: 1.652369737625122
Average Height: 1.5309635400772095
PCA Components for 95% Variance: 7
Optimal K-Means Clusters: 4
DBSCAN Clusters: 263
DBSCAN Noise Points: 835

Summary saved to: C:\Users\94675\Desktop\group_project_pro\logs\data_exploration_
summary.txt
```